

Detector-coupled Monte Carlo estimate of background and unresolved-line sensitivity for a balloon-borne 511 keV Laue-lens TES telescope

Authors omitted for review

Abstract

The Galactic 511 keV annihilation signal is dominated by diffuse bulge and disk emission, while compact or point-like contributions remain poorly constrained. We present a detector-coupled Monte Carlo estimate for a balloon-borne pointed telescope combining a 10 m Ge(111) Laue lens with a transition-edge sensor microcalorimeter focal plane. The signal chain transports focused 511 keV photons through the Be aperture into the detector/cryostat mass model. The background chain transports prompt atmospheric particles and photons and production-position-sampled delayed activation decays through the same model. Prompt, delayed, and signal streams are analyzed with a common 510.58–511.42 keV window, CsI anticoincidence veto, and topology/field-of-view consistency selection. For a reference point-source flux of $10^{-4} \text{ ph cm}^{-2} \text{ s}^{-1}$, the selected day-15 signal and background rates are $1.18 \times 10^{-3} \text{ s}^{-1}$ and $4.81 \times 10^{-2} \text{ s}^{-1}$, respectively; prompt atmospheric events account for 91.7% of the selected background. Folding these rates through the adopted synthetic 20 d trajectory gives $Z_{20d} = 7.04$ and a corresponding 3σ flux threshold of $4.26 \times 10^{-5} \text{ ph cm}^{-2} \text{ s}^{-1}$. The result is a selected-rate flight-sensitivity estimate for a spectrally unresolved compact 511 keV source under the reference geometry and response model. Prompt atmospheric events dominate the selected background, while the delayed component is smaller but materially structured, identifying shielding, side-entry collimation, local materials, and activation validation as the main next engineering and analysis targets.

Keywords Gamma-ray instrumentation, 511 keV annihilation line, transition-edge sensors, Laue lens, balloon-borne telescope, activation background

In this work we develop an end-to-end detector-coupled Monte Carlo chain for this concept and use it to estimate the in-flight background and point-source performance at balloon altitude and to provide directional guidance for design optimization. The chain takes the EXPACS/PARMA atmospheric cosmic-ray field as input and, at balloon altitude, transports the prompt particles while tracking their activation products, yielding both the prompt atmospheric background and the delayed activation background that builds up over the flight; the focused 511 keV reference signal is provided by a 10 m Ge(111) Laue-optics model. The signal and both backgrounds are injected into the same representative balloon-borne TES focal-plane mass model (a dilution-refrigerator cold end, passive shielding, and an active anticoincidence shield), processed with a common narrow energy window, anticoincidence veto, and topology/field-of-view selection, and folded along a reference flight trajectory to give the selected signal rate, the background composition, and the counting significance. The selected background is dominated by the prompt atmospheric component, so the shield geometry, side-entry collimation, and shield material are the most direct directions for design optimization; delayed activation enters the same evaluation as a secondary component, whereas upstream optics-hardware background remains outside the primary detector-background budget. This work thereby estimates the in-flight background and detection performance of the balloon platform and indicates directions for design optimization.

0.1 Science and instrument requirements

The simulated instrument is aimed at 511 keV point-source or unresolved-central-excess components from the Galactic-center region that have not been detected by previous instruments. Existing INTEGRAL/SPI measurements have established a 511 keV sky dominated by diffuse bulge and disk emission and limit point-like 511 keV sources superposed on that diffuse emission to fluxes of order $10^{-4} \text{ ph cm}^{-2} \text{ s}^{-1}$ [Siegert et al.(2016), Yoneda et al.(2025)]; an independent INTEGRAL/IBIS compact-source search also found no Galactic 511 keV point sources and gives an upper limit of about $1.6 \times 10^{-4} \text{ ph cm}^{-2} \text{ s}^{-1}$ for a ~ 10 Ms Galactic-center exposure [De Cesare(2011)]. For a representative 511 keV atmospheric transmission of $T \simeq 0.739$ at balloon altitude, this limit corresponds to an above-payload photon-flux scale of about $1.2 \times 10^{-4} \text{ ph cm}^{-2} \text{ s}^{-1}$. The

simulation therefore adopts $F_0 = 10^{-4} \text{ ph cm}^{-2} \text{ s}^{-1}$ as the reference point-source photon-flux scale. Compared with the wide-field spectroscopic and coded-mask searches performed by INTEGRAL, the modeled payload adds Laue focusing to improve the point-source response and uses a high-energy-resolution, multilayer thick-absorber TES microcalorimeter array as the focal-plane detector to sharpen the 511 keV line measurement [Shirazi et al.(2023), Caputo et al.(2017)].

This work first considers a balloon payload as an early test platform, with the expectation that the approach can later be extended to a satellite observing platform.

This goal imposes four requirements on the instrument and the simulation: first, the telescope must concentrate 511 keV photons onto a small focal plane, so that a larger solid angle can be covered without sacrificing collecting area. Second, the focal-plane detector must provide sub-keV energy response sufficient to resolve detailed line structure. Third, the detector and shield model must suppress atmospheric, activation, and side-entry backgrounds. Fourth, the simulation chain must treat prompt background, delayed activation, focused signal, vetoes, event topology, and mission-time variation consistently, so that signal and background are evaluated under the same detector selection.

The remainder of this paper is organized as follows. Section 1 describes the detector and cryostat mass model; Section 2 the simulation framework—the workflow, optics coupling, background sources, event selection, and mission-time fold; Section 3 the selected-rate and sensitivity results together with the normalization, convergence, and boundary checks; and Section 4 the discussion and conclusions.

1 Detector and cryostat mass model

The detector mass model is an engineering Geant4/MEGAlib proxy mass model that contains a balloon-borne multilayer TES array mounted at the cold end of a compact dilution refrigerator (DR) [Shirazi et al.(2023), Gottardi & Nagayoshi(2021)]. The detector model also contains typical DR cryostat structures, active shielding, passive shielding, magnetic shielding, side-entry windows, collimation structures, copper thermal structures, and other cryogenic service-mass proxies. Figure 1 shows two-dimensional views of the mass model used in this paper.

The TES detector concept uses an absorber-coupled AlMn film structure developed in our laboratory; the advantages of AlMn alloy include transition-temperature tunability through composition and annealing, suitability for array fabrication, and relatively low magnetic-field sensitivity [Xie et al.(2026), Vavagiakis et al.(2017)]. In the present concept, the TES thermometer film is an AlMn film of order 200 nm thickness, whereas the γ -ray absorber is a millimetre-scale metal body. The absorber uses $1.5 \times 1.5 \times 3.0 \text{ mm}^3$ Ta pixels; the high atomic number of Ta increases the interaction probability for 511 keV photons, and a 3.0 mm thickness corresponds to about a 49% single-layer detection efficiency at 511 keV, so the six-layer array can provide near-unity detection efficiency. In the current geometry, the center-to-center spacing between adjacent Ta array layers is 12 mm; this spacing is designed with the efficiency of the subsequent Compton-based kinematic-track veto mechanism in mind. The choice of Ta also considers its low-temperature low-specific-heat properties as a superconductor (this paper uses an order-of-magnitude estimate of $10^{-2} \text{ J K}^{-1} \text{ m}^{-3}$) and a relatively high photoelectric-to-Compton interaction ratio; based on NIST XCOM cross sections near 511 keV, this ratio is about 0.731 [Hubbell & Seltzer]. Based on laboratory experience and by substituting the current Ta absorber heat capacity of 67.5 pJ/K, the simulation adopts $\Delta E_{\text{FWHM}} \simeq 420 \text{ eV}$ at 511 keV as the current simulated energy resolution.

Because the TES film and Ta absorber differ greatly in size and mass, and because the TES film itself does not contribute significant 511 keV interactions, the transport geometry uses the Ta absorber to represent the main sensitive volume of a TES pixel. The explicit sensitive volume has six layers, 376 Ta absorber pixels per layer, and 2256 pixels in total, with 1.5 mm length and width and 3 mm thickness. A Si substrate proxy is placed below each Ta pixel, using Si to represent the SiO₂-SiN_x-Si substrate because this simulation does not model detailed thermal conduction. TES-film, wiring, and readout details are not built layer by layer as explicit interaction volumes; instead, they enter the model through the assumed energy response and service-mass proxies. In the current manuscript, the 511 keV line window is $W_{511} = 511 \text{ keV} \pm 420 \text{ eV}$, i.e. 510.58–511.42 keV.

The TES array is mounted on the required copper support and thermal-link structures. Copper provides the decisive thermal-conduction path; a copper thermal finger exits through the bottom hole of the sample box and connects to the 50 mK cold plate. The near-field TES sample box serves as both magnetic shielding and




Figure 1: Two-dimensional overview of the detector/cryostat proxy mass model, showing the X-Z view of the DR tower, detector bay, active shield, and 511 keV side-entry beam path.

local passive shielding. The current mainline geometry consists of a 2 mm Nb inner layer and a 2 mm MuMetal outer layer, with MuMetal represented in the material table as Ni:Fe = 4:1. The sample-box entrance window is covered by a thin Al window; the rear end cap retains the hole for the cold finger. A W–Nb double-layer sample box is treated as a comparison passive-shielding + magnetic-shielding material design; the X-ray window is opened on the side of the cryostat, with a 300 μm thick Al window embedded in the thermal-shield shell inside each cryostat layer, a 25 μm thick Be window placed in the vacuum shell, and an x -thick W collimator placed outside the Be window, with the collimator opening strictly aligned to the TES pixels. The side window is strictly defined through Boolean operations. The mass model also introduces scintillator active shielding, wrapped around the outer side wall, bottom, and top of the cryostat. Segmented CsI active shielding is used, with a radial side-shield thickness of 4 cm; the bottom and top shields are included in the geometry as axial end volumes. CsI volumes enter the simulation as transport and activation materials, and their deposited energies are read in post-processing to apply the active veto; the counting threshold is 50 keV. BGO active shielding is treated as a material-comparison case; the BGO threshold is 70 keV. Considering the physical layout of the DR and the side-entry optical path, the cryostat instrument frame is rotated as a whole by 45 degrees, so that the local side window is tilted 45° upward in the zenith-frame source coordinates.

2 Simulation framework

Detector-coupled Monte Carlo transport through detailed mass models is a mature method for predicting the response and background of gamma-ray spectrometers [Sturmer et al.(2003)]. This section describes the full end-to-end simulation chain from source generation to flight-performance assessment. The focusing optics and detector system start from independent mass models: the optics-branch mass model takes the target point source and records focused photons reaching the focal-plane plane as the focused-source input to the detector branch; the detector branch then directly transports, in the detector/cryostat mass model of Section 1, prompt atmospheric particles, delayed decays generated from activation products according to recorded production positions, and the focused sources described above. The focused source, prompt source, and delayed source are simulated stepwise and normalized onto one time axis through Poisson-statistics-based time-axis sampling and renormalization, then the background is further suppressed by scintillator-based active veto and Compton-kinematics-based veto. The chain subsequently enters detailed background-analysis statistics, analysis of signal and background intensity variation caused by mission flight-time trajectory changes, and the overall flight-performance assessment and design-optimization recommendations. Because the focused signal and the two background streams all pass through the same response and veto definitions before any background decomposition or sensitivity statement, the subsequent results rest on a consistent event selection.

Figure 2 expands this chain by computational layer: geometry definition, prompt and activation-accumulation transport, delayed-source construction, focused-optics transport, detector-coupled signal replay, veto mechanisms, detailed background analysis, time/flight-trajectory variation, and flight-performance assessment with design-optimization recommendations. The following subsections describe each layer's inputs, operations, and outputs.

The above simulation is for a specific balloon-flight moment, longitude and latitude, and altitude. Although for an observation time such as one hour equivalent to several Monte Carlo events, this position can be regarded as constant. However, for a long-time flight on the scale of days, the altitude, longitude, and latitude obviously change significantly. This causes changes in the magnetic-rigidity cutoff for instantaneous cosmic rays and changes in the atmospheric transmission for the signal, and the activation background itself also changes with the incident cosmic-ray flux and different stages of the decay chain. This reminds us that the background-signal at a reference moment should not be linearly extrapolated in time to evaluate the time-folded flight performance over the whole working time. Therefore, we introduce a reasonable analytical flight-trajectory-time curve to evaluate the performance over the whole working cycle, and simulate multiple points on the curve to verify the rationality of the curve.




Figure 2: Simulation chain used in this work. Blue nodes are the prompt and delayed background streams, green nodes the focused-signal optics, and grey nodes the shared geometry, veto mechanisms, detailed background analysis, time/trajectory variation, and final flight-performance assessment. The focused signal and all background streams are evaluated under the same detector response and veto definitions before background decomposition, mission-time variation, and design-optimization judgments are made.

2.1 Laue optics and detector-coupled signal replay

This simulation adopts a laue-lens system based on mosaic Ge(111) crystals [Barriere et al.(2009), Barriere et al.(2011)]. Through optical design, it realizes a 10 m focal length, a 511 keV single-ring reference response, and an effective area of 20.1 cm² at 511 keV; the effective area is calculated as

$$A_{\text{eff}}(E) = A_{\text{inc}} \frac{N_{\text{fp}}(E)}{N_{\text{inc}}(E)} \quad (1)$$

where A_{inc} is the Monte Carlo incident sampling area, $N_{\text{inc}}(E)$ is the number of incident photons, and $N_{\text{fp}}(E)$ is the number of photons satisfying the focusing/focal-plane interception condition. The mosaic system is adopted to make the Monte Carlo calculation convenient while satisfying the expected focusing performance; other candidate schemes are also possible for actual engineering adaptation.

In this simulation, the laue optical branch gives the focusing-performance simulation and transports the focused signal through the Ge crystal mass.

For this purpose, the diffraction probability of photons by the crystal is no longer superposed onto photons as an efficiency parameter, but is defined as a discrete process at the Geant4 application layer [Agostinelli et al.(2003), Allison et al.(2016)]. The process takes the diffraction probability $R(|\Delta\theta|)$ according to the angular deviation of the current photon relative to the crystal plane, and converts it into a finite equivalent diffraction mean free path λ_{diff} , so that the accumulated diffraction probability when the photon crosses a crystal path ℓ reproduces this value: $1 - \exp(-\ell/\lambda_{\text{diff}}) = R(|\Delta\theta|)$, namely $\lambda_{\text{diff}} = -\ell / \ln[1 - R(|\Delta\theta|)]$; to avoid double counting with standard photoelectric absorption, R is first normalized by the unabsorbed fraction ($R/(1 - A)$) and then converted into a mean free path. After this mean free path is returned to Geant4, the diffraction process and the standard electromagnetic processes (photoelectric absorption and Compton scattering) each sample interaction lengths in the same competition framework, and the shortest one decides

which process occurs first, rather than forcibly intercepting photons at the crystal boundary. Photons not selected to undergo diffraction still continue transport according to the standard electromagnetic physics, and may be absorbed, scattered, or directly transmitted.

The diffraction probability is preferentially given by interpolation in an external XOP/CRYSTAL rocking-curve map as a function of angular deviation: this curve is calculated offline once in the preprocessing stage by the CRYSTAL/diff_pat module of the XOP toolkit (called through xoppylib, with DABAX crystallographic data as input) [Sanchez del Rio & Dejus(2011)], and gives the Ge(111) reflectivity/transmissivity/absorption curves as a function of angular deviation; during transport, the angular deviation is calculated in real time for every step inside the crystal, and the current diffraction probability is obtained by interpolation on that curve, namely “offline library tabulation, probability used by angle during transport”. The mosaic orientation is implemented by applying a Gaussian angular perturbation to the ideal Bragg-plane normal, with the perturbation width given by the 30 arcsec mosaic FWHM; after diffraction occurs, the outgoing direction is mirror-reflected according to the perturbed plane normal, and the photon energy is unchanged. This split in which the “process class only handles the transport interface, and the independent model class handles probability and direction decisions” follows the idea of existing crystal Bragg-reflection extensions in Geant4 [Guan et al.(2023), Reiazi et al.(2025)], and extends it to the energy band needed here.

Here the reference signal is a monochromatic 511 keV point source whose flux scale is set by the INTEGRAL/SPI upper limit on a compact 511 keV source superposed on the Galactic-centre diffuse emission ($F_0 = 10^{-4} \text{ ph cm}^{-2} \text{ s}^{-1}$ above the atmosphere, Section 0.1) [Siebert et al.(2016), De Cesare(2011)], rather than by any single observed object. It is constructed as a MEGAlib far-field source [Zoglauer et al.(2006)], with the detailed construction given in Section 2.2.3. The signal is transported through the optics branch and injected into the detector mass model as a parallel-beam source; prompt cosmic-ray background and activation decays are transported directly in the detector/cryostat mass model. Optics-mass prompt and delayed background are not included in the primary detector-background budget.

2.2 Source construction

After the focusing optics have been defined in Section 2.1, the optical input and the subsequent detector input require three source layers: the prompt atmospheric particle field incident on the payload, the delayed radioactive population accumulated in the instrument by that particle field, and the reference signal point source. These three layers are constructed separately, rather than running transport by directly invoking MEGAlib’s buildup function for them at the same time, in order to reduce computational cost.

2.2.1 Prompt atmospheric source

The prompt atmospheric field is derived from EXPACS/PARMA. PARMA3.0 is an analytical model parameterized from PHITS atmospheric-shower calculations and benchmarked against terrestrial cosmic-ray measurements, and EXPACS is its public tabulation and software interface [Sato(2015), Sato(2016)]. PHITS supplies the underlying particle-transport physics: primary cosmic rays interact with air nuclei, the secondary hadrons, leptons, photons, and neutrons are transported through the atmospheric column, and the resulting fluences are parameterized by altitude, geomagnetic cutoff, solar modulation, energy, and direction. In this work EXPACS/PARMA provides the differential incident flux as a function of particle species, kinetic energy, zenith angle, geographic position, altitude, cutoff rigidity, and solar activity, and is used only to construct the external radiation fields; detector response, vetoing, topology selection, and mission-time counting are applied downstream. The reference source definitions use latitude 34° , longitude 100° , altitude 38 km, vertical cutoff rigidity $R_c = 11.6 \text{ GV}$, and the 2025-08-31 solar condition corresponding to $W = 118.3$.

The prompt source includes photons, neutrons, electrons, positrons, protons, alpha particles, and negative and positive muons. For each particle family the EXPACS/PARMA differential spectrum is converted into a MEGAlib far-field area source covering the full zenith range $\theta = 0^\circ\text{--}180^\circ$ and full azimuth range $\phi = 0^\circ\text{--}360^\circ$. The zenith dependence is represented by 20 differential equal- μ angular bins, where $\mu = \cos \theta$ and each bin subtends $\Delta\Omega = 0.6283185307 \text{ sr}$: the first ten bins cover down-going particles, the last ten up-going or albedo-like particles. Both hemispheres are retained in the prompt transport and in the activation-production transport,




Figure 3: EXPACS/PARMA full-sphere atmospheric flux field used to generate the prompt and activation-production transport inputs. Left: energy- and solid-angle-integrated fluxes for the eight transported particle families, separated into down-going and up-going hemispheres. Right: bin-integrated flux in the 20 equal- μ differential zenith bins used by the MEGAlib far-field source. The plot is generated from the source-definition inputs for latitude 34°, longitude 100°, altitude 38 km, $R_c = 11.6$ GV, and $W = 118.3$.

so upward albedo particles are not discarded at source construction. Figure 3 shows the resulting full-sphere flux field.

The integrated source-plane fluxes are approximately 4.80, 0.462, 0.199, 0.117, 0.112, 0.0115, 0.0050, and 0.0056 cm⁻² s⁻¹ for γ , n , e^- , e^+ , p , α , μ^- , and μ^+ , respectively; the non-photon families together carry about 0.91 cm⁻² s⁻¹, or 16% of the total. If Monte Carlo counts were allocated strictly in proportion to flux, the detector and activation statistics of these minority families would be poor, so the non-photon families are generated with replicated samples and de-weighted in the rate normalization. All particles are then transported through the detector/cryostat mass model with Geant4/MEGAlib [Agostinelli et al.(2003), Allison et al.(2016), Zoglauer et al.(2006)]; the prompt run and the activation-production run each generated 25,210,216 incident primaries. Each particle species is normalized with its own transported exposure,

$$r_{\text{event},j} = \left(\sum_i \mathcal{T}_{ij} \right)^{-1}, \quad (2)$$

where $\mathcal{T}_{ij}$ is the transported exposure of subsample i for species j . This per-species normalization is required because the photon and non-photon streams have different sampling and exposure structures; the normalization check enforces $r_{\text{event},j} \sum_i \mathcal{T}_{ij} = 1$ for every species, so over-sampling changes the Monte Carlo variance but not the expected physical rate.

2.2.2 Delayed activation source

For the delayed simulation, a separate activation-production stream is run by enabling BUILDUP during the prompt-background transport. This activation-production stream is not counted as an immediate detector background. Instead, it records the isotope, material volume, and production position of each radioactive product. These records are converted through the production-decay accumulation equation into a radioactive population fixed to a chosen reference time during the flight, with ground-state half-lives checked against NUBASE2020 [Kondev et al.(2021)]. For isotope k ,

$$A_k(t_{\text{flight}}) = P_k [1 - \exp(-\lambda_k t_{\text{flight}})], \quad \lambda_k = \frac{\ln 2}{t_{1/2,k}}, \quad (3)$$

where P_k is the production rate inferred from the activation-production transport and $t_{\text{flight}} = 15$ d for the reference state. The fixed delayed total activity used by the reference delayed-source construction is 141.83 Bq.

The position information does not come from the aggregate isotope store alone. The standard store records volume names and isotope counts; the production-position construction adds a custom recording hook in the MEGAlib stepping-action layer. Whenever a delayed isotope is stored during activation buildup, the hook writes a production-position record at the same moment, including the logical volume name, pre-step position (x, y, z) , isotope identifier $1000Z + A$, excitation energy, time, process, and track-ancestry information. The fixed day-15 population is then matched to these records by volume, isotope, and excitation state, so the delayed source can use the simulated production coordinates rather than only volume-integrated activities. The current Mass_model_511 source-level inventory is dominated by CsI-shield activation and ^{128}I , with smaller contributions from Cu, Al, W, Mg, and Cs isotopes. A Mass-specific appendix position-distribution figure remains to be regenerated. It is a source-construction diagnostic, not a detector-count map; detector response is applied only after the delayed decays are transported through the common geometry and event selection.

Whether the production positions are worth preserving can be tested from the same table. Delayed decays are emitted locally, so downstream attenuation, active-shield veto probability, and event topology all depend on where the isotope was produced, and different incident families—with their different energy-loss, stopping, and capture histories—need not deposit their activation in the same materials. To quantify this, the weighted production rows are grouped by incident particle family and material category, and the activity-normalized category distributions of two families a and b are compared with the total-variation distance

$$D_{\text{TV}}(a, b) = \frac{1}{2} \sum_c |f_{a,c} - f_{b,c}|, \quad (4)$$

where $f_{a,c}$ is the fraction of delayed activity from family a produced in material category c . In the fixed day-15 table, neutron-induced products carry about 133.67 Bq, μ^- -induced products 7.63 Bq, proton-induced products 0.434 Bq, and α -induced products 0.0825 Bq, with smaller μ^+ , e^+ , and γ terms. The neutron-induced activity remains dominated by the CsI shield. The corresponding Mass_model_511 material-category fractions have not yet been regenerated. A volume-integrated or axisymmetric radial source could preserve the total activity, but it would erase these incident-family, material, nuclide, and coordinate correlations; the production-position construction keeps them in the source before detector transport.

The delayed-decay source is therefore emitted from sampled production positions rather than from an axisymmetric distribution. Let row j in the weighted production table have activity weight w_j , with total activity $A = \sum_j w_j$. During construction, M production positions are drawn with replacement using probability $p_j = w_j/A$, and each sampled point source is assigned the same flux A/M . The expected activity assigned to row j is

$$\mathbb{E} \left[N_j \frac{A}{M} \right] = M p_j \frac{A}{M} = w_j, \quad (5)$$

so the estimator preserves the total activity and is unbiased for the weighted spatial distribution. Here M is the number of equal-activity sampled point sources; it is a sampling size, not the number of physical production rows. The reference delayed source preserves the fixed total activity at the source-definition level and is replayed through the same detector/cryostat geometry as the prompt stream. The nominal calculation transports 10^6 delayed decays; the corresponding Cosima live time is only the transport record used to generate the decay sequence and is not used as the physical balloon exposure time.

Table 1 summarizes the source layers used in the current detector-coupled calculation. It is placed in the source-model section because this section defines incident particle fields and delayed radioactive populations; detector vetoes and topology selections are applied only in the subsequent selection section.

This production-position representation does not create one sampled point source for every production row. It is a statistical sampling approximation whose adequacy remains to be evaluated at the selected-rate level. Source-list checks verify activity conservation and source serialization; selected-rate behavior remains to be tested after detector transport and event selection.

2.2.3 Reference point-source construction

The focused signal stream is defined by a monochromatic 511 keV reference point source. Its flux scale is set not by any single observed object but by the INTEGRAL/SPI upper limit on a compact 511 keV source

Table 1: Background-source layers used by the current detector-coupled chain. The table defines source construction only; anticoincidence, topology selection, and the line window are applied later in Section 2.3.

Layer	Physical role	Current implementation
Prompt atmospheric transport	Direct detector hits from atmospheric secondaries	EXPACS/PARMA full-sphere source definitions at 34° latitude, 100° longitude, 38 km altitude, $R_c = 11.6$ GV, and $W = 118.3$; γ , n , $e^\pm$, p , α , and $\mu^\pm$ over 20 equal- μ differential zenith bins including both down-going and up-going hemispheres; 25,210,216 generated primaries.
Activation production	Radioisotope production and production-position history	Same atmospheric field transported as a BUILDUP activation-production stream, again retaining both hemispheres and all eight particle families; 25,210,216 generated primaries.
Fixed delayed population	Radioactive activity at the adopted day-15 state after continuous exposure	Production-decay accumulation with NUBASE-checked ground-state half-lives; total fixed activity 141.83 Bq.
Production-position decay transport	Delayed decay particles emitted from sampled production positions	Equal-activity sampled point sources from the weighted production table; reference transport uses 10^6 delayed decays per source sampling.

superposed on the Galactic-centre diffuse emission, $F_0 = 10^{-4} \text{ ph cm}^{-2} \text{ s}^{-1}$ above the atmosphere (Section 0.1) [Siegert et al.(2016), De Cesare(2011)].

Because such a source lies at astronomical distance, it is modelled as an ideal on-axis point: in the optics run it is a monochromatic 511 keV far-field beam whose photons arrive parallel to the optical axis and are sampled uniformly over the incident aperture A_{inc} of Eq. (1). No finite angular size is assigned to the source itself; the angular structure that reaches the focal plane is produced entirely by the lens—the 30 arcsec mosaic spread, the XOP/CRYSTAL rocking curve, and the $f = 10$ m focusing geometry—so that the replayed beam converges to within about 0.4° of the axis and forms a focused spot rather than a single ray. Three independent seeds of 5×10^4 primaries each give 1.5×10^5 transported photons, of which 37,256 are diffracted and cross the focal plane; 37,194 of these (99.8%) fall within the $r = 1.898$ cm Be entrance window, with focused-spot radii $r_{90} = 1.03$ cm and $r_{99} = 1.25$ cm. The within-window crossings form the event list that is replayed through the Be aperture into the detector/cryostat model with the recorded positions and directions.

The normalization ties this replay to the reference flux. At the detector entrance the focused stream carries the rate $F_0 A_{\text{eff}}(511 \text{ keV}) T$, where $A_{\text{eff}}(511 \text{ keV}) = 20.1 \text{ cm}^2$ is the on-axis effective area and $T = 0.739$ is the reference 511 keV atmospheric transmission at the day-15 trajectory anchor; numerically this is $1.48 \times 10^{-3} \text{ s}^{-1}$ at the reference flux. Equivalently, the generating run corresponds to an observation time $N_{\text{fp}}/[F_0 A_{\text{eff}} T]$, and each replayed event carries the inverse of that exposure as its rate weight; in the intermediate cut-flow bookkeeping of Section 2.3 the same sample appears as a unit replay before this scale is applied. After the common detector response and event selection, the surviving signal rate is $1.185 \times 10^{-3} \text{ s}^{-1}$, corresponding to a selected effective area of $S_{W_{511}}/F_0 = 11.85 \text{ cm}^2$. Because the optics run contains only source photons, the replayed stream is signal-pure; all backgrounds enter through the streams of Section 2.2.

The effective area quoted above is the on-axis value, representing best-case pointing. A point source offset from the optical axis, or residual attitude jitter during flight, reduces and shifts the focused response through the finite angular acceptance (field of view) of the Laue lens [Barriere et al.(2009), Weidenspointner et al.(2006)]; this off-axis dependence is treated as a pointing systematic on the signal effective area rather than as a change to the on-axis reference construction. The current 45° side-entry response is a north-sky/near-zenith reference, not a simulated low-elevation Galactic-centre observation.




Figure 4: Rate normalization on a common coincidence time axis. (a) Full-band detector event rates (928, 97.1, and 0.987 s^{-1}). (b) Independent Poisson draws of the three streams are merged and time-ordered; color marks stream identity and hatched bands mark each event’s temporal footprint. (c) Events with $\Delta t \leq \tau = 1 \mu\text{s}$ form one candidate group c ; the box gives the reference-exposure accidental bookkeeping. Tick positions are schematic.

2.3 Detector event selection and veto definitions

The detector-level selection chain places the focused-signal, prompt atmospheric, and delayed-activation streams on one common time axis and applies two veto layers in sequence. Because both veto layers are defined through time coincidence, the time axis and its rate normalization are established first, and the vetoes are defined on top of it.

2.3.1 Time axis and rate normalization

Each stream enters with a per-record rate weight r_{kj} from the source normalization of Section 2.2, where k labels the stream and j labels a record in that stream’s event catalogue; the stream total rate is $R_k = \sum_j r_{kj}$. To place three independently normalized catalogues on one physical time axis, the streams are combined by a Poisson superposition: for stream k the number of instances is drawn as

$$N_k \sim \text{Pois}(R_k T), \quad (6)$$

records are selected with probabilities r_{kj}/R_k with replacement, each is assigned a uniform random time in the reference window $[0, T]$, and all streams are merged and time-ordered into a single detector time axis carrying the three physical rates. Event instances separated by no more than $\tau = 1 \mu\text{s}$ are grouped into a coincidence candidate group c . In the baseline (Figure 4) the full-band detector rates superposed on this axis are about 928 s^{-1} (prompt), 97.1 s^{-1} (delayed activation), and 0.987 s^{-1} (focused signal), so the coincidence axis is strongly prompt-dominated; over the reference exposure the 7,215,233 event instances yield 1275 mixed-stream coincidences, giving an accidental live factor $\simeq 0.9990$ (the line window and shield threshold below then act on the resulting candidate groups). The normalized summed-energy spectrum of the three streams is shown in Figure 5: the window is line-dominated, with a starved continuum.

This post-processing superposition, rather than the native trigger/veto flags present in the geometry description, defines the rate for two reasons. First, both veto layers are defined by time coincidence: the active anticoincidence gate compares the TES and active-shield deposits within the same τ window, and the topology reconstruction must first decide which hits belong to the same incident photon; both require a common time axis carrying all three streams at their physical rates. Second, the common post-processing definition forces the prompt, delayed, and focused-signal streams through the same coincidence window τ and the same shield threshold, so that production-stage trigger flags do not silently become the rate definition. The superposition is licensed by a stationarity assumption: within the adopted reference exposure each stream rate is treated as constant. The reference window is short compared with the timescales over which the physical rates drift—the




Figure 5: Normalized summed-energy spectrum of the three streams on the common axis (raw stage). The focused signal, prompt, and delayed streams all concentrate at the 511 keV line while the continuum is starved. Shape-diagnostic re-derivation on the surviving `Mass_model_511` catalogues; absolute selected rates are authoritative in Table 2.

atmospheric column and the activation inventory evolve on the mission-day scale, whereas on the microsecond-to-second scale relevant to coincidence the three rates are effectively stationary. A constant-rate arrival process is a homogeneous Poisson process, and the superposition of independent homogeneous Poisson processes is again Poisson with rate $\sum_k R_k$, which is exactly what licenses drawing $N_k \sim \text{Pois}(R_k T)$ and merging. The mission-time significance accordingly uses the expectation values of these rates rather than adding a fresh random draw to every mission bin; the Poisson realization only validates the grouping and accidental-coincidence bookkeeping.

2.3.2 Active anticoincidence gate

The first layer is an active anticoincidence gate on this common time axis. For each candidate group c the TES and active-shield energy sums are

$$E_{\text{TES}}^{(c)} = \sum_{j \in c} E_{\text{TES},j}, \quad E_{\text{sh}}^{(c)} = \sum_{j \in c} E_{\text{sh},j}. \quad (7)$$

The group enters the hard line window if $E_{\text{TES}}^{(c)}$ falls in

$$W_{511} = [510.58, 511.42] \text{ keV} \quad (511 \text{ keV} \pm 420 \text{ eV}), \quad (8)$$

which is the current `Mass_model_511` selection window, retaining most of the line response while suppressing continuum background. The group passes the active shield if

$$E_{\text{sh}}^{(c)} < E_{\text{veto}}, \quad E_{\text{veto}} = 50 \text{ keV}, \quad (9)$$

where for the baseline CsI configuration the shield sum is read from the CsI active-shield volumes in post-processing. This layer is the dominant source of rate-level background suppression. In the baseline it reduces the prompt background from 153 to 71 events ($0.1180 \rightarrow 0.0482 \text{ s}^{-1}$, 40.8% surviving) and the delayed activation from 43 to 28 events (65.1%), while retaining all 30305 focused-signal events (zero shield deposit, 100%); see Table 2. Its effect on the background spectrum is shown in Figure 6.

The quoted flux thresholds apply to a line whose intrinsic width, after detector response, remains small compared with this window; a substantially broadened source line would require convolving the source line profile with the response and re-optimizing the window.




Figure 6: Effect of the CsI anticoincidence gate on the prompt+delayed background spectrum (in-window pre-veto vs after the shield gate). Shape-diagnostic re-derivation on the surviving `Mass_model_511` catalogues.

2.3.3 Compton topology and field-of-view consistency

The second layer applies a Compton-kinematic topology and field-of-view consistency test to the TES hits that survive the active-shield gate; its physical basis is Compton-camera cone reconstruction [Zoglauer et al.(2006)]. An incident photon of energy E_0 Compton-scatters at a first interaction point $\mathbf{r}_1$, depositing E_1 , and the scattered photon of energy $E_{sc} = E_0 - E_1$ propagates to a second interaction point $\mathbf{r}_2$; the scattering angle follows from

$$\cos \theta_C = 1 - m_e c^2 \left(\frac{1}{E_{sc}} - \frac{1}{E_1 + E_{sc}} \right). \quad (10)$$

Given the scattered-photon direction $\hat{\mathbf{u}}_{12} = (\mathbf{r}_2 - \mathbf{r}_1)/|\mathbf{r}_2 - \mathbf{r}_1|$, the source direction is constrained to a reconstruction cone with apex $\mathbf{r}_1$, axis $-\hat{\mathbf{u}}_{12}$, and half-angle θ_C , whose projection on the sky is the event circle. This work does not perform imaging; it only tests whether this first-scatter cone intersects the side Be-window disk (radius 1.898 cm), i.e. whether the reconstructed arrival direction is consistent with entry through the aperture.

Because the interaction position within a single TES pixel ($1.5 \times 1.5 \times 3.0$ mm) is unresolved, each hit is represented not by its centroid alone but by $8 + 1 = 9$ representative points—the eight pixel corners and the centroid—carrying the sub-pixel position uncertainty explicitly into the geometric reconstruction. A two-hit event in a given scatter order therefore enumerates $9 \times 9 = 81$ first/second position combinations and constructs one reconstruction cone per combination (each cone sampled at 24 azimuths, propagated to the Be-window plane, and tested against the disk by both its sampled points and the connecting segments); the order is deemed field-of-view consistent if any of these 81 trajectories yields a cone that intersects the disk. This is a deliberately permissive consistency test: an event is retained if any sub-pixel interaction position is consistent with entry through the aperture. The layer consequently retains about 95% of the in-window multi-hit events and removes only about 10–14% of the background, consistent with its role as a loose consistency veto rather than an imaging reconstruction.

The three event classes are handled accordingly. Single-pixel events have no scatter topology and are retained as calorimetric line candidates. A two-hit event provides two positions and energies but no internal information to order the interactions, so both scatter orders are enumerated over the 81 trajectories above, and the event is retained if either order is field-of-view consistent and rejected only if neither order admits any kinematically valid cone. Three-hit and higher-multiplicity events contain internal scatter vertices and can be ordered from geometric-kinematic consistency: at each internal vertex the scattering angle can be obtained both from the deposited energies through the Compton formula and from the angle between successive flight directions, and the two should agree for the correct order. All orders are enumerated (up to six TES hits; higher multiplicities




Figure 7: (a) Background spectrum through the two veto layers (raw; after anticoincidence; after anticoincidence+Compton/FoV). (b) Hit multiplicity of the selected candidate groups per stream. Shape-diagnostic re-derivation on the surviving `Mass_model_511` catalogues; absolute rates are authoritative in Table 2.

are retained without ordering), the best order minimizes the Compton sequence residual

$$Q = \sum_{i=1}^{n-2} (\cos \theta_{C,i} - \hat{\mathbf{u}}_{i-1,i} \cdot \hat{\mathbf{u}}_{i,i+1})^2 \quad (11)$$

(ties broken by the longest first-scatter lever arm), and the same 81-trajectory cone-disk test is applied to its first two points. Events with no valid order are classified as unreconstructed. In the baseline, events that pass the active shield are sorted by the topology test into single-site, field-of-view-consistent reconstructed multi-hit, and vetoed classes: the prompt stream comprises 43 single-site, 22 reconstructed multi-hit, and 6 vetoed events (71 total); the delayed stream 20, 8, and 0 (28 total); and the focused signal 17589 single-site, 12098 reconstructed multi-hit, and 618 vetoed events (30305 total), giving 65, 28, and 29687 selected events, respectively. Reclassifying the current `Mass_model_511` detector geometry in the 480–550 keV broad band by TES hit multiplicity (Figure 7; shape diagnostic only; absolute selected rates remain those in Table 2) resolves the selected focused-signal groups into 17598 single-site, 9794 two-hit, and 2340 ≥ 3 -hit events (with 370 two-hit and 262 ≥ 3 -hit events Compton-vetoed); multi-hit counts in the prompt and delayed streams remain small (two-/ ≥ 3 -hit: 25/7 and 12/3).

Unreconstructed events are retained in the baseline selected-rate definition but are not counted as field-of-view passing, so the line-window result does not gain sensitivity by discarding ambiguous reconstruction failures. This policy is also not a conservative lower bound on the net sensitivity: a single-site-only alternative that vetoes all multi-hit candidates would retain 59.2% of the focused signal (17589/29687) and about 66.6% of the selected background in the present W_{511} sample, reducing the simple $S/\sqrt{B}$ figure to about 0.73 of the baseline value. The baseline selection therefore treats ambiguous multi-hit events as a reconstruction-model boundary to be validated, not as a source of sensitivity gain from aggressive event rejection.

To connect this selection chain explicitly to the finite Monte Carlo sample, Table 2 gives the current event-level W_{511} cut-flow obtained by lightweight replay of the simulated event catalogues. The table is not a new Geant4/MEGAlib transport. It replays the current event catalogues with the same TES energy window, CsI veto, and Compton-kinematic/field-of-view consistency checks; uncertainties are Monte Carlo quadrature sums over event weights. The current event-selection records first fix the energy window and then report raw, active-shield, and topology stages, so the raw stage should be read as the pre-veto sample inside the line window, not as an all-energy incident sample.

As a retained reference-geometry cross-check that has not been rerun for `Mass_model_511`, the focused-signal replay stream was reconstructed with the MEGAlib `revan` Compton event reconstructor (37,065 events processed, reading the existing simulation without re-running transport) [Zoglauer et al.(2006)]. The class-level results agree with the custom chain: in-window single-pixel events number 17,695 (`revan`) versus 17,663 (custom

Table 2: event-level W_{511} line-window cut-flow.

Stream	Stage	Events	Rate / s^{-1}	Statistical uncertainty / s^{-1}
Prompt background	Line-window pre-veto	153	1.1802×10^{-1}	1.25×10^{-2}
Prompt background	After CsI veto	71	4.8160×10^{-2}	5.72×10^{-3}
Prompt background	Topology/FOV selected	65	4.4089×10^{-2}	5.47×10^{-3}
Delayed activation	Line-window pre-veto	43	6.1192×10^{-3}	9.33×10^{-4}
Delayed activation	After CsI veto	28	3.9846×10^{-3}	7.53×10^{-4}
Delayed activation	Topology/FOV selected	28	3.9846×10^{-3}	7.53×10^{-4}
Focused signal	Line-window pre-veto	30305	8.1478×10^{-1}	4.68×10^{-3}
Focused signal	After CsI veto	30305	8.1478×10^{-1}	4.68×10^{-3}
Focused signal	Topology/FOV selected	29687	7.9817×10^{-1}	4.63×10^{-3}

The focused-signal rows are unit-replay rates used to show veto survival behavior; the point-source signal rate in Table 3 is scaled separately by the Laue effective area, reference flux, and atmospheric transmission.

chain), a 0.2% difference, while the 3.2% difference in the in-window total (31,295 versus 30,323) follows from reapplying a per-hit Gaussian energy response and a trigger threshold that feed the partial-absorption shoulder below the window into it, predominantly in the multi-hit class, whereas the custom chain takes the window on unnoised summed energies. For two-hit ordering the cross-check gives two quantitative results. First, the fraction of two-hit events with at least one kinematically valid order is 100.0% (10,421/10,421), as expected analytically—for a fully absorbed 511 keV two-hit event the scattered-photon energy cannot fall below the $E_0/3$ Compton edge, so at least one order is always valid—confirming the premise for retaining two-hit events. Second, the best order chosen by the Klein–Nishina cross section with a kinematic quality factor agrees with the order that points more nearly toward the source only 78.4% of the time, so committing to a single estimated order (as an imaging chain would) would use the wrong Compton cone for about 22% of two-hit events, whereas enumerating both orders and accepting either field-of-view-consistent one avoids this misassignment. The cross-check also yields a quantitative bound absent from the earlier chain: the median $|\text{ARM}|$ of in-window Compton events relative to the true beam direction is about 24° (two-hit 23.6° ; three-hit and higher 26.3° ; only about 32% have $|\text{ARM}| < 10^\circ$), driven by the millimetre-scale scatter lever arms and pixel quantization; this explains why the topology/field-of-view layer removes only about 10–14% of the background and cannot be sharpened into an imaging-level discriminant. Finally, the fraction of unorderable events is very small (61 of 37,065, i.e. 0.16%, in revan; one in the custom chain), showing that the unreconstructed class is a small boundary population rather than a hidden pillar of the sensitivity.

81 points from 0 to 20 d with 0.25 d spacing; endpoint points carry half-widths in the count integration. The trajectory is a synthetic reference profile, not balloon telemetry. For time-grid point t_i in days,

$$h_i = h_0 + \Delta h \sin \left[\frac{2\pi(t_i - 15)}{5} \right], \quad (12)$$

$$\varphi_i = \varphi_0 + \Delta\varphi \sin \left[\frac{2\pi(t_i - 15)}{7} \right], \quad (13)$$

$$\ell_i = \ell_0 + \Delta\ell \sin \left[\frac{2\pi(t_i - 15)}{9} \right], \quad (14)$$

where $h_0 = 38.75$ km, $\Delta h = 2.5$ km, $\varphi_0 = 34^\circ$, $\Delta\varphi = 0.25^\circ$, $\ell_0 = 100^\circ$, and $\Delta\ell = 0.25^\circ$. The resulting ranges are 36.25–41.25 km in altitude, 33.75° – 34.25° in latitude, and 99.75° – 100.25° in longitude.

The atmospheric depth and the 511 keV transmission scale are

$$X(h_i) = X_{38} \exp \left[-\frac{h_i - 38 \text{ km}}{H} \right], \quad X_{38} = 3.87 \text{ g cm}^{-2}, \quad H = 6.8 \text{ km}, \quad (15)$$

$$T_i = \exp[-\mu_{\text{eff}} X(h_i)], \quad \mu_{\text{eff}} = -\frac{\ln T_{15}}{X(h_0)}, \quad (16)$$

with $T_{15} = 0.739$. Thus the focused-science scale is $f_{\text{atm},i} = T_i/T_{15}$. The prompt atmospheric source scale is

$$R_{c,i} = 11.0 - 0.08(\varphi_i - \varphi_0) + 0.03(\ell_i - \ell_0), \quad (17)$$

$$f_{p,i} = \exp\left[\frac{X(h_i) - X(h_0)}{30 \text{ g cm}^{-2}}\right] \left(\frac{11.0}{R_{c,i}}\right)^{0.2}. \quad (18)$$

The trajectory rigidity model is centered on a rounded value, and $f_{p,i}$ enters only through the day-15–anchored ratio $(11.0/R_{c,i})^{0.2}$, so this central value affects the mission-time variation but not the absolute prompt rate, which is fixed by the $R_c = 11.6$ GV source definition. The same scalar production driver is used for delayed activation before isotope-by-isotope decay buildup is solved. For a delayed-production element a with decay constant λ_a and fixed day-15 activity $A_{a,15}$,

$$P_{a,15} = \frac{A_{a,15}}{1 - \exp(-\lambda_a t_{15})}, \quad (19)$$

$$\frac{dN_a}{dt} = P_{a,15} f_p(t) - \lambda_a N_a, \quad A_a(t) = \lambda_a N_a(t). \quad (20)$$

The numerical fold anchors the solution back to $A_{a,15}$ at day 15 and sums the activity over the nuclide/material/position production elements.

This delayed-activity treatment assumes that the normalized activation-product distribution is unchanged over the small trajectory range, and that the trajectory changes only the scalar production driver. The verification described here checks the internal implementation of that model but does not prove physical invariance of the production distribution. If dedicated activation transports at the trajectory center and extrema show material or nuclide redistribution, the scalar driver should be replaced by interpolation over production vectors,

$$P_a(t_i) = \sum_g w_g(h_i, \varphi_i, \ell_i, R_{c,i}) P_a^{(g)}, \quad \sum_g w_g = 1, \quad (21)$$

followed by the same decay ODE.

Implementation checks recomputed the trajectory formula at the center point and at the two altitude extrema. The reference atmospheric transmission is $T = 0.739$ at the day-15 anchor, and the high- and low-altitude probes give $T = 0.811$ and 0.646 , respectively. Reconstructing the W_{511} prompt, delayed, background, and signal rates from the day-15 rates and these scales reproduces the mission-fold summary. Reweighting the particle families with local PARMA spectra at the same trajectory probes shifts the normalized nuclide–volume activity distribution by less than 6×10^{-3} in total-variation distance. This supports using a scalar delayed production/activity fold for the present small trajectory range at the particle-family reweighting layer. It is not a full proof of distribution invariance: same-particle energy/angle-dependent activation cross sections and position-transport shifts would require dedicated low-, reference-, and high-altitude activation-production transports.

In the mission-time fold, the selected day-15 rates are folded over the trajectory grid as

$$B_i = L_i (B_{p,15} f_{p,i} + B_{d,15} f_{d,i}), \quad (22)$$

$$S_i = L_i S_{15} f_{\text{atm},i}, \quad (23)$$

$$L_i = \exp[-(R_{p,i} + R_{d,i})\tau], \quad (24)$$

where $f_{p,i}$ is the prompt atmospheric scale, $f_{d,i}$ is the delayed-activity scale, $f_{\text{atm},i}$ is the 511 keV atmospheric transmission scale, and L_i is the accidental coincidence live factor. Here $R_{p,i}$ and $R_{d,i}$ are the prompt and delayed event rates entering the coincidence grouping at that mission point. The factor L_i assumes a one-sided τ gate; a symmetric $\pm\tau$ condition would instead give $\exp[-2(R_{p,i} + R_{d,i})\tau]$, but for the present rates both forms remain very close to unity, so the folded counts are insensitive to this choice.

The time-dependent significance is computed as a cumulative counting statistic. For a reference flux F_0 , the cumulative significance is

$$Z(t_m) = \frac{N_s(t_m)}{\sqrt{N_b(t_m)}}, \quad (25)$$

where

$$N_s(t_m) = \sum_{i \leq m} S_i \Delta t_i, \quad N_b(t_m) = \sum_{i \leq m} B_i \Delta t_i \quad (26)$$

are the time-integrated selected signal and background counts after mission-time scaling and live-factor corrections. For comparison among analysis variants, the 3σ flux threshold is scaled as

$$F_{3\sigma}(t_m) = F_0 \frac{3}{Z(t_m)}. \quad (27)$$

This Gaussian counting convention is used consistently for the line-window results and secondary comparisons, but it is a diagnostic selected-rate metric, not a final flight discovery likelihood. For the present 20 d reference case, the signal-to-background ratio is about 2.4%; a 1% Gaussian uncertainty on the background normalization alone would reduce the approximate significance from about 7.0 to about 2.3. A flight analysis would therefore require off-source or control-band background constraints and nuisance-parameter profiling. For the primary 20 d line-window result, the folded counts are about 8.3×10^4 background counts and 2.0×10^3 signal counts at F_0 , so this approximation is used only in the high-count selected-rate regime. The isolated upstream Ge-proxy result with zero selected events is reported separately with a Poisson zero-count upper limit. The detector-selection summary also records a day-15 direct expectation. The primary result does not use that constant-rate direct value; it uses the mission-time fold and the final cumulative significance. The distinction is small for the present line-window case, but keeping it explicit is necessary for comparing analysis configurations.

3 Results

The recovered log fragments around the original detailed results tables are structurally mixed. This compile-ready copy therefore keeps the confirmed primary selected-rate numbers in a compact recovery table and leaves the original detailed table reconstruction to manual editing.

Table 3: Recovered primary unresolved-line selected-rate result.

Quantity	Value	Note
Day-15 selected background	$4.81 \times 10^{-2} \text{ s}^{-1}$	Prompt + delayed
Day-15 selected signal at F_0	$1.18 \times 10^{-3} \text{ s}^{-1}$	$F_0 = 10^{-4} \text{ ph cm}^{-2} \text{ s}^{-1}$
Prompt background fraction	91.7%	Dominant component
20 d diagnostic Z	7.04	Counting metric
$F_{3\sigma}(20 \text{ d})$	$4.26 \times 10^{-5} \text{ ph cm}^{-2} \text{ s}^{-1}$	Reference model only

Recovery gap note

The original recovered log fragments become structurally inconsistent after this point. The broad-window diagnostic paragraph and several nearby transition lines are omitted from this compile-ready reading copy; see the raw partial source and recovery manifest.

4 Discussion and conclusions

The main result is that the reference production-position delayed-source calculation gives a 20 d diagnostic 3σ point-source threshold of $4.26 \times 10^{-5} \text{ ph cm}^{-2} \text{ s}^{-1}$ for an unresolved line in the modeled focused TES line spectrometer. This threshold is obtained without using the focused-spot spatial analysis as the primary result. It follows from three combined effects: optical concentration, a sub-keV energy window, and a detector-coupled veto/topology selection that suppresses prompt and delayed backgrounds in the same selection space as the signal.

This calculation should be compared with existing 511 keV instruments as a pointed, narrow-line, compact-source estimate, not as an all-sky imaging forecast. INTEGRAL/SPI and COSI provide the reference measurements for mapping the diffuse 511 keV sky and the broader Galactic diffuse emission [Siegert et al.(2016), Kierans et al.(2020), Siegert et al.(2020), Yoneda et al.(2025), Karwin et al.(2023)]. Peer-reviewed COSI calibration and simulation studies provide the more relevant technical comparison for detector response, event reconstruction, and balloon-background validation [Beechert et al.(2022), Sleator et al.(2019), Gallego et al.(2025a), Gallego et al.(2025b), Ciabatonni et al.(2025)]. The TES Laue-lens concept considered here is complementary: it trades sky coverage for a concentrated beam and a sub-keV line window. The relevant comparison class is compact 511 keV targets, unresolved-central-excess searches, and transient follow-up observations with known or constrained source positions. The same focused aperture is not a substitute for measuring the full diffuse bulge and disk luminosity, because it samples only a small solid angle.

The production-position delayed-source treatment addresses a specific approximation in radial-profile compression. Although delayed activation contributes only about 8.3% of the selected line-window background in the present calculation, it is spatially and materially structured. Preserving sampled production positions prevents artificial azimuthal or radial redistribution from becoming an unquantified systematic. The current Mass_model_511 result uses one full-stat production-position sampling; a same-statistics multi-seed selected-rate convergence check remains open.

The background composition indicates the first engineering parameters to test, rather than a completed optimization. The selected line-window rate remains dominated by prompt atmospheric streams, so shield geometry, side-entry collimation, material near the focal plane, topology selection, and multi-pixel reconstruction are the most direct handles on the present background. Delayed activation is smaller but remains spatially and materially structured, so source-position preservation is retained in the rate calculation. The dominance of iodine activation in the CsI shield also motivates a future shield-material trade study.

4.1 Limitations and future validation

The selected-rate result has five immediate boundaries. First, the detector/cryostat model is an engineering proxy, not a final flight CAD model. Second, quantitative active-shield rejection is applied in the common detector-level post-processing selection rather than by production-time trigger/veto filtering. Third, the mission-time calculation is a rate-level fold, not a new particle transport at each trajectory time bin. Fourth, the custom Compton/field-of-view veto is a detector-level topology and kinematic consistency check, not yet a full Revan/Mimrec reconstruction, and the retained reconstruction-failure class is a baseline acceptance choice rather than a conservative sensitivity bound. Fifth, the event-level analysis has made the cut-flow, broad-window spectra, veto behavior, and activation selected-event decomposition traceable, but selected delayed state/metastable labels, the merged four-sampling per-event decomposition, post-veto energy-window staging, multi-pixel residual distributions, and veto-threshold stability scans are still not closed by the current outputs.

More generally, the tabulated 3σ flux thresholds are statistical selected-rate thresholds for the reference source, mass, and response model. They do not include a systematic-uncertainty envelope for atmospheric-

flux normalization, activation cross sections or physics-list choices, material and geometry tolerances, detector energy-response uncertainty, veto-threshold calibration, reconstruction-efficiency systematics, or background-estimation nuisance parameters. A publication-level flight sensitivity should replace the diagnostic $S/\sqrt{B}$ metric with a spatial-spectral likelihood, following standard profile-likelihood practice [Cowan et al.(2011)].

The retained upstream Laue-lens calculation should be read as a boundary calculation for an equal-volume/equal-mass active-Ge proxy, not as a finalized optics-background budget. It completes full-particle prompt and activation-production transport for the proxy, and the isolated Ge-proxy delayed component gives zero selected W_{511} events in the retained transport. The prompt self-background contribution is still not part of the primary detector-background budget because it needs source separation or subtraction from the ordinary detector/cryostat prompt background, and explicit lens support hardware is not yet modeled.

Spatial analyses such as r_{90} and multi-annulus likelihood estimates are diagnostic cross-checks, but they are not yet a finalized nuisance-profile likelihood. The quoted sensitivity therefore remains the line-window selected-rate result, not a spatially optimized likelihood result.

Before this rate study can support a flight-performance statement, four validation blocks remain: a Re-van/Mimrec or equivalent reconstruction cross-check of the retained multi-hit subset, including residual distributions and threshold stability; optics self-background transport with explicit lens tiles and support hardware at final statistics; a final cryostat/DR geometry and detector-response update, including TES saturation and pile-up constraints; and a spatial-spectral likelihood that profiles background nuisance parameters. The event-level tables added here should be read as manuscript support evidence and audit entry points before those validations, not as a complete systematic-uncertainty envelope.

4.2 Conclusions

We have presented a detector-coupled background and sensitivity study for a balloon-borne pointed 511 keV TES Laue-lens telescope. The reference calculation couples a Ge(111) Laue optical response to a TES focal-plane mass model, prompt atmospheric transport, delayed activation, active-shield vetoing, topology/field-of-view selection, mission-time rate folding, and counting-significance evaluation. In the primary 510.58–511.42 keV unresolved-line window, the day-15 selected rates are $B \simeq 4.81 \times 10^{-2} \text{ s}^{-1}$ and $S \simeq 1.18 \times 10^{-3} \text{ s}^{-1}$ at $10^{-4} \text{ ph cm}^{-2} \text{ s}^{-1}$; the mission-time fold gives $Z_{20d} \simeq 7.04$ and $F_{3\sigma}(20d) \simeq 4.26 \times 10^{-5} \text{ ph cm}^{-2} \text{ s}^{-1}$. The result provides a detector-coupled selected-rate basis for continued evaluation of focused TES spectroscopy as a complementary 511 keV technique for compact-source and unresolved-central-excess searches, while identifying prompt atmospheric background suppression, delayed-activation systematic validation, and reconstruction validation as the main next validation tasks.

Data and software availability

Before publication, a versioned public repository or institutional archive should provide the geometry and source definitions, random seeds, analysis and figure-generation scripts, LaTeX source, and derived validation products needed to reproduce the tables and figures in this paper. This compile-ready recovery copy is not the original final manuscript.

Declarations

Funding To be completed before journal submission.

Competing interests To be completed before journal submission.

Author contributions To be completed before journal submission.

Ethics approval and consent to participate Not applicable.

References

[Prantzos et al.(2011)] N. Prantzos, C. Boehm, A.M. Bykov, R. Diehl, K. Ferriere, N. Guessoum, P. Jean, J. Knoedlseder, A. Marcowith, I.V. Moskalenko, A. Strong, G. Weidenspointner, The 511

keV emission from positron annihilation in the Galaxy, *Rev. Mod. Phys.* 83 (2011) 1001–1056, doi:10.1103/RevModPhys.83.1001.

- [Jean et al.(2006)] P. Jean, J. Knoedlseder, V. Lonjou, M. Allain, J.-P. Roques, G.K. Skinner, B.J. Teegarden, G. Vedrenne, P. von Ballmoos, B. Cordier, P. Caraveo, R. Diehl, P. Durouchoux, P. Leleux, R. Schanne, G. Weidenspointner, Spectral analysis of the Galactic e^+e^- annihilation emission, *Astron. Astrophys.* 445 (2006) 579–589, doi:10.1051/0004-6361:20053765.
- [Siegert et al.(2016)] T. Siegert, R. Diehl, G. Khachatryan, M.G.H. Krause, F. Guglielmetti, J. Greiner, A.W. Strong, X. Zhang, Gamma-ray spectroscopy of positron annihilation in the Milky Way, *Astron. Astrophys.* 586 (2016) A84, doi:10.1051/0004-6361/201527510.
- [Siegert et al.(2019)] T. Siegert, R. Diehl, G. Khachatryan, M.G.H. Krause, F. Guglielmetti, J. Greiner, A.W. Strong, X. Zhang, Constraints on positron annihilation kinematics in the inner Galaxy, *Astron. Astrophys.* 627 (2019) A126, doi:10.1051/0004-6361/201833856.
- [Yoneda et al.(2025)] H. Yoneda, T. Siegert, S. Mittal, Imaging the positron annihilation line with 20 years of INTEGRAL/SPI observations, arXiv:2509.01066, 2025.
- [Kierans et al.(2020)] C.A. Kierans, S.E. Boggs, A. Zoglauer, A.W. Lowell, C. Sleator, J. Beechert, T.J. Brandt, P. Jean, H. Lazar, J. Roberts, T. Siegert, J.A. Tomsick, P. von Ballmoos, Detection of the 511 keV Galactic positron annihilation line with COSI, *Astrophys. J.* 895 (2020) 44, doi:10.3847/1538-4357/ab89a9.
- [Siegert et al.(2020)] T. Siegert, S.E. Boggs, J.A. Tomsick, A.C. Zoglauer, C.A. Kierans, C.C. Sleator, J. Beechert, T.J. Brandt, P. Jean, H. Lazar, A.W. Lowell, J.M. Roberts, P. von Ballmoos, Imaging the 511 keV positron annihilation sky with COSI, *Astrophys. J.* 897 (2020) 45, doi:10.3847/1538-4357/ab9607.
- [Knoedlseder et al.(2005)] J. Knoedlseder, P. Jean, V. Lonjou, G. Weidenspointner, N. Guessoum, R. Diehl, W. Gillard, M.J. Harris, G.K. Skinner, P. von Ballmoos, G. Vedrenne, P. Caraveo, B. Cordier, V. Schoenfelder, B.J. Teegarden, The all-sky distribution of 511 keV electron-positron annihilation emission, *Astron. Astrophys.* 441 (2005) 513–532, doi:10.1051/0004-6361:20042063.
- [De Cesare(2011)] G. De Cesare, Searching for the 511 keV annihilation line from galactic compact objects with the IBIS gamma ray telescope, *Astron. Astrophys.* 531 (2011) A56, doi:10.1051/0004-6361/201116516.
- [Shirazi et al.(2023)] F. Shirazi, E. Gau, M.A. Hossen, D. Becker, D. Schmidt, D. Swetz, D. Bennett, D. Braun, F. Kislat, J. Gard, J. Mates, J. Weber, N.R. Cavero, S. Chun, L. Lisalda, A. West, B. Dev, F. Ferrer, R. Bose, J. Ullom, H. Krawczynski, The 511-CAM Mission: A pointed 511 keV gamma-ray telescope with a focal plane detector made of stacked transition edge sensor microcalorimeter arrays, *J. Astron. Telesc. Instrum. Syst.* 9 (2023) 024006, doi:10.1117/1.JATIS.9.2.024006.
- [Barriere et al.(2009)] N.M. Barriere, L. Natalucci, N. Abrosimov, P. von Ballmoos, P. Bastie, P. Courtois, M. Jentschel, J. Knoedlseder, J. Rousselle, P. Ubertini, Soft gamma-ray optics: new Laue lens design and performance estimates, arXiv:0910.0488, 2009.
- [Barriere et al.(2011)] N.M. Barriere, J.A. Tomsick, S.E. Boggs, A. Lowell, P. von Ballmoos, Developing a second generation Laue lens prototype: high reflectivity crystals and accurate assembly, arXiv:1111.6700, 2011.
- [Weidenspointner et al.(2006)] G. Weidenspointner, C.B. Wunderer, N. Barriere, A. Zoglauer, P. von Ballmoos, Monte Carlo study of detector concepts for the MAX Laue Lens gamma-ray telescope, arXiv:astro-ph/0603152, 2006.
- [Sanchez del Rio & Dejus(2011)] M. Sanchez del Rio, R.J. Dejus, XOP v2.4: recent developments of the x-ray optics software toolkit, *Proc. SPIE* 8141 (2011) 814115, doi:10.1117/12.893911.

- [Guan et al.(2023)] F. Guan, M. Asai, D.A. Bartkoski, M. Kleckner, Z. Harel, M. Salehpour, Adding the X-ray Bragg reflection physical process in crystal to the Geant4 Monte Carlo simulation toolkit, part I: reflection from a crystal slab, *Precision Radiation Oncology* 7(1) (2023) 59–66.
- [Reiazi et al.(2025)] R. Reiazi, D. Lee, D. Bartkoski, M. Kleckner, M. Salehpour, G4BraggReflection for accurate modeling of Bragg reflection in perfect and mosaic crystals, *J. Phys. D: Appl. Phys.* 58 (2025) 295102, doi:10.1088/1361-6463/aded1d.
- [Siegert et al.(2016)] T. Siegert, R. Diehl, J. Greiner, M.G.H. Krause, A.M. Beloborodov, M. Cadolle Bel, F. Guglielmetti, J. Rodriguez, A.W. Strong, X. Zhang, Positron annihilation signatures associated with the outburst of the microquasar V404 Cygni, *Nature* 531 (2016) 341–343, doi:10.1038/nature16978.
- [Gottardi & Nagayoshi(2021)] L. Gottardi, K. Nagayoshi, A review of X-ray microcalorimeters based on superconducting transition edge sensors for astrophysics and particle physics, *Applied Sciences* 11 (2021) 3793, doi:10.3390/app11093793.
- [Xie et al.(2026)] L. Xie, Y. Zhang, Z. Li, Z. Liu, S. Shu, J. Zhou, X. Li, H. Li, H. Gao, Y. Gu, X. Lu, Y. Zhao, C. Liu, Beyond one-thousandth energy resolution with an AlMn TES detector, arXiv:2602.11728, 2026.
- [Vavagiakis et al.(2017)] E.M. Vavagiakis, S.W. Henderson, K. Zheng, H.-M. Cho, N.F. Cothard, B. Dober, S.M. Duff, P.A. Gallardo, G. Hilton, J. Hubmayr, K.D. Irwin, B.J. Koopman, D. Li, F. Nati, M.D. Niemack, C.D. Reintsema, S. Simon, J.R. Stevens, A. Suzuki, B. Westbrook, Magnetic sensitivity of AlMn TESes and shielding considerations for next-generation CMB surveys, arXiv:1710.08456, 2017.
- [Hubbell & Seltzer] J.H. Hubbell, S.M. Seltzer, XCOM: Photon Cross Section Database, NIST Standard Reference Database 8 (XGAM), National Institute of Standards and Technology, doi:10.18434/T48G6X.
- [Zhang et al.(2022)] S. Zhang, J.-K. Xia, T. Sun, et al., Transition edge sensor based detector: from X-ray to gamma-ray, arXiv:2204.00010, 2022.
- [Gehrels(1985)] N. Gehrels, Instrumental background in balloon-borne gamma-ray spectrometers and techniques for its reduction, *Nucl. Instrum. Methods Phys. Res. A* 239 (1985) 324–349, doi:10.1016/0168-9002(85)90732-6.
- [Sato(2015)] T. Sato, Analytical model for estimating terrestrial cosmic ray fluxes nearly anytime and anywhere in the world: extension of PARMA/EXPACS, *PLoS ONE* 10 (2015) e0144679, doi:10.1371/journal.pone.0144679.
- [Sato(2016)] T. Sato, Analytical model for estimating the zenith angle dependence of terrestrial cosmic ray fluxes, *PLoS ONE* 11 (2016) e0160390, doi:10.1371/journal.pone.0160390.
- [Kondev et al.(2021)] F.G. Kondev, M. Wang, W.J. Huang, S. Naimi, G. Audi, The NUBASE2020 evaluation of nuclear physics properties, *Chinese Phys. C* 45 (2021) 030001, doi:10.1088/1674-1137/abddae.
- [Sturmer et al.(2003)] S.J. Sturmer, C.R. Shrader, G. Weidenspointner, B.J. Teegarden, D. Attie, B. Cordier, R. Diehl, C. Ferguson, P. Jean, A. von Kienlin, P. Paul, F. Sanchez, S. Schanne, P. Sizun, G. Skinner, C.B. Wunderer, Monte Carlo simulations and generation of the SPI response, *Astron. Astrophys.* 411 (2003) L81–L84, doi:10.1051/0004-6361:20031171.
- [Agostinelli et al.(2003)] S. Agostinelli, J. Allison, K. Amako, et al., Geant4—a simulation toolkit, *Nucl. Instrum. Methods Phys. Res. A* 506 (2003) 250–303, doi:10.1016/S0168-9002(03)01368-8.
- [Allison et al.(2016)] J. Allison, K. Amako, J. Apostolakis, et al., Recent developments in Geant4, *Nucl. Instrum. Methods Phys. Res. A* 835 (2016) 186–225, doi:10.1016/j.nima.2016.06.125.
- [Zoglauer et al.(2006)] A. Zoglauer, R. Andritschke, F. Schopper, MEGAlib: The Medium Energy Gamma-ray Astronomy Library, *New Astron. Rev.* 50 (2006) 629–632, doi:10.1016/j.newar.2006.06.049.

- [Caputo et al.(2017)] R. Caputo, F. Kislat, J. Racusin, on behalf of the AMEGO Team, AMEGO: Simulations of the instrument performance, Proc. 35th International Cosmic Ray Conference, PoS(ICRC2017) 783 (2018), doi:10.22323/1.301.0783.
- [Karwin et al.(2023)] C.M. Karwin, T. Siebert, J. Beechert, J.A. Tomsick, T.A. Porter, M. Negro, C. Kierans, M. Ajello, I. Martinez Castellanos, A. Shih, A. Zoglauer, S.E. Boggs, Probing the Galactic diffuse continuum emission with COSI, arXiv:2310.12206, 2023.
- [Gallego et al.(2025a)] S. Gallego, U. Oberlack, J. Lommler, C.M. Karwin, A. Zoglauer, P. Jean, P. von Ballmoos, C. Kierans, C.C. Sleator, J.A. Tomsick, S.E. Boggs, Bottom-up background simulations of the 2016 COSI balloon flight, arXiv:2503.02493, 2025.
- [Gallego et al.(2025b)] S. Gallego, U. Oberlack, J. Lommler, C.M. Karwin, F. Fenu, V. Fioretti, A. Zoglauer, et al., Pre-flight background estimates for COSI, arXiv:2510.25304, 2025.
- [Beechert et al.(2022)] J. Beechert, H. Lazar, S.E. Boggs, T.J. Brandt, Y.-C. Chang, C.-Y. Chu, H. Gulick, C. Kierans, A. Lowell, N. Pellegrini, J.M. Roberts, T. Siebert, C.C. Sleator, J.A. Tomsick, A. Zoglauer, Calibrations of the Compton Spectrometer and Imager, Nucl. Instrum. Methods Phys. Res. A 1031 (2022) 166510, doi:10.1016/j.nima.2022.166510.
- [Sleator et al.(2019)] C.C. Sleator, A. Zoglauer, A.W. Lowell, C.A. Kierans, S.E. Boggs, J.A. Tomsick, H. Lazar, J. Beechert, T.J. Brandt, P. von Ballmoos, Benchmarking simulations of the Compton Spectrometer and Imager with calibrations, Nucl. Instrum. Methods Phys. Res. A 946 (2019) 162643, doi:10.1016/j.nima.2019.162643.
- [Ciabattini et al.(2025)] A. Ciabattini, V. Fioretti, J.A. Tomsick, A. Zoglauer, P. Patel, L. Mitchell, A. Bulgarelli, P. Jean, G. Panebianco, N. Parmiggiani, C. Vignali, P. von Ballmoos, E. Wulf, Benchmarking of Geant4 simulations for the COSI anticoincidence system, Experimental Astronomy (2025), doi:10.1007/s10686-025-10019-7.
- [Cowan et al.(2011)] G. Cowan, K. Cranmer, E. Gross, O. Vitells, Asymptotic formulae for likelihood-based tests of new physics, Eur. Phys. J. C 71 (2011) 1554, doi:10.1140/epjc/s10052-011-1554-0.